

An introduction to



Jianguo (Jeff) Xia, Assistant Professor

Institute of Parasitology, and Department of Animal Science

McGill University

June 1, 2017

<http://www.xialab.ca>

Objective & outline

- To get comfortable with R!


1. Why R?
2. Review basic R commands
3. Review basic R data types
 - Vector, matrix, list, data frame
4. Review basic R graphics
 - Scatter plot, box plot, histogram
5. Write a simple R script!

Survey

- What is your experience in R?
 - a. Never used at all
 - b. Used in my work but very basic
 - c. Using it rather regularly in my work
- Do you have experience in any other language?
 - Java, JavaScript, C, C++, ...
 - Perl, Python, Ruby, ...

The Shift of Landscape (Google Trend)

- Decline of Bioinformatics



Rise of Data Science

























<https://opensource.com/business/14/12/r-open-source-language-data-science>

Embracing Data Science

- Bioinformatics is now considered part of growing the field of Data Science
- Data science is an interdisciplinary field about processes and systems to extract knowledge or insights from data in various forms, either structured or unstructured
- Big data is embraced by all big players:
 - Academics: NIH, EBI
 - Industry: Google, Amazon, IBM, Microsoft, Facebook

The 2016 Top Programming Languages

C is No. 1, but big data is still the big winner

Language Rank	Types	Spectrum Ranking
1. C	  	100.0
2. Java	  	98.1
3. Python	 	98.0
4. C++	  	95.9
5. R		87.9
6. C#	  	86.7
7. PHP		82.8
8. JavaScript	 	82.2
9. Ruby	 	74.5
10. Go	 	71.9

<http://spectrum.ieee.org/computing/software/the-2016-top-programming-languages>



What is R

- A programming language
- A computing environment for:
 - Statistics
 - Machine learning
 - Graphics
- Open source
 - Main language in teaching statistics
 - Main statistics language at Google
 - Main language in statistical analysis of “omics” data
 - Microarray, RNAseq, metagenomics, metabolomics

Working with R

- Command line interface
- A text editor
 - Linux: nedit, gedit
 - Mac: TextMate, TextWrangler
 - Windows: Notepad ++
- R Studio
 - <https://www.rstudio.com/>



Syntax

- The **\$** indicate terminal command line
- The **>** is for R command line
- The arrow **<-** is the assignment operator in R
 - Equal sign **=** is now supported but **<-** is preferred
- The **#** is for comments in script
 - Message for yourself or other developer
 - Ignored by R during execution

```
> weight.a <- 10
> weight.a
[1] 10
# get sum of numbers
> sum(c(2, 6, 8))
```

Interactive R command line

1. Open a terminal
2. Create a new, empty working directory
3. Navigate to the directory
4. Start R
5. Play around
 - `getwd()` # current position
 - `dir()` # show files in current directory
 - `setwd(...)` # move to new dir
6. Exit R: `q()`

```
$ mkdir ReviewR
$ cd ReviewR
$ R
> getwd()
[1] "/Users/xia/ReviewR"
> dir()
character(0)
> print ("hello world!")
[1] "hello world!"
> q()
```

A Powerful Calculator

- R can directly evaluate math expressions.
- Special symbols:
 - **NaN**: Not a Number
 - **Inf**: infinite number
 - **NA**: Not Available or missing value)
 - **NULL**: undefined value
 - **TRUE**: logical value true
 - **FALSE**: logical value false

```
> 2+2
[1] 4
> 2^8
[1] 256
> exp(-2)
[1] 0.1353353
> log(-10)
[1] NaN
> 1/0
[1] Inf
```

Finding Help: ?

```
> ?apply
```

```
apply                                package:base                                R Documentation
```

```
Apply Functions Over Array Margins
```

```
Description:
```

```
Returns a vector or array or list of values obtained by applying a  
function to margins of an array or matrix.
```

```
Usage:
```

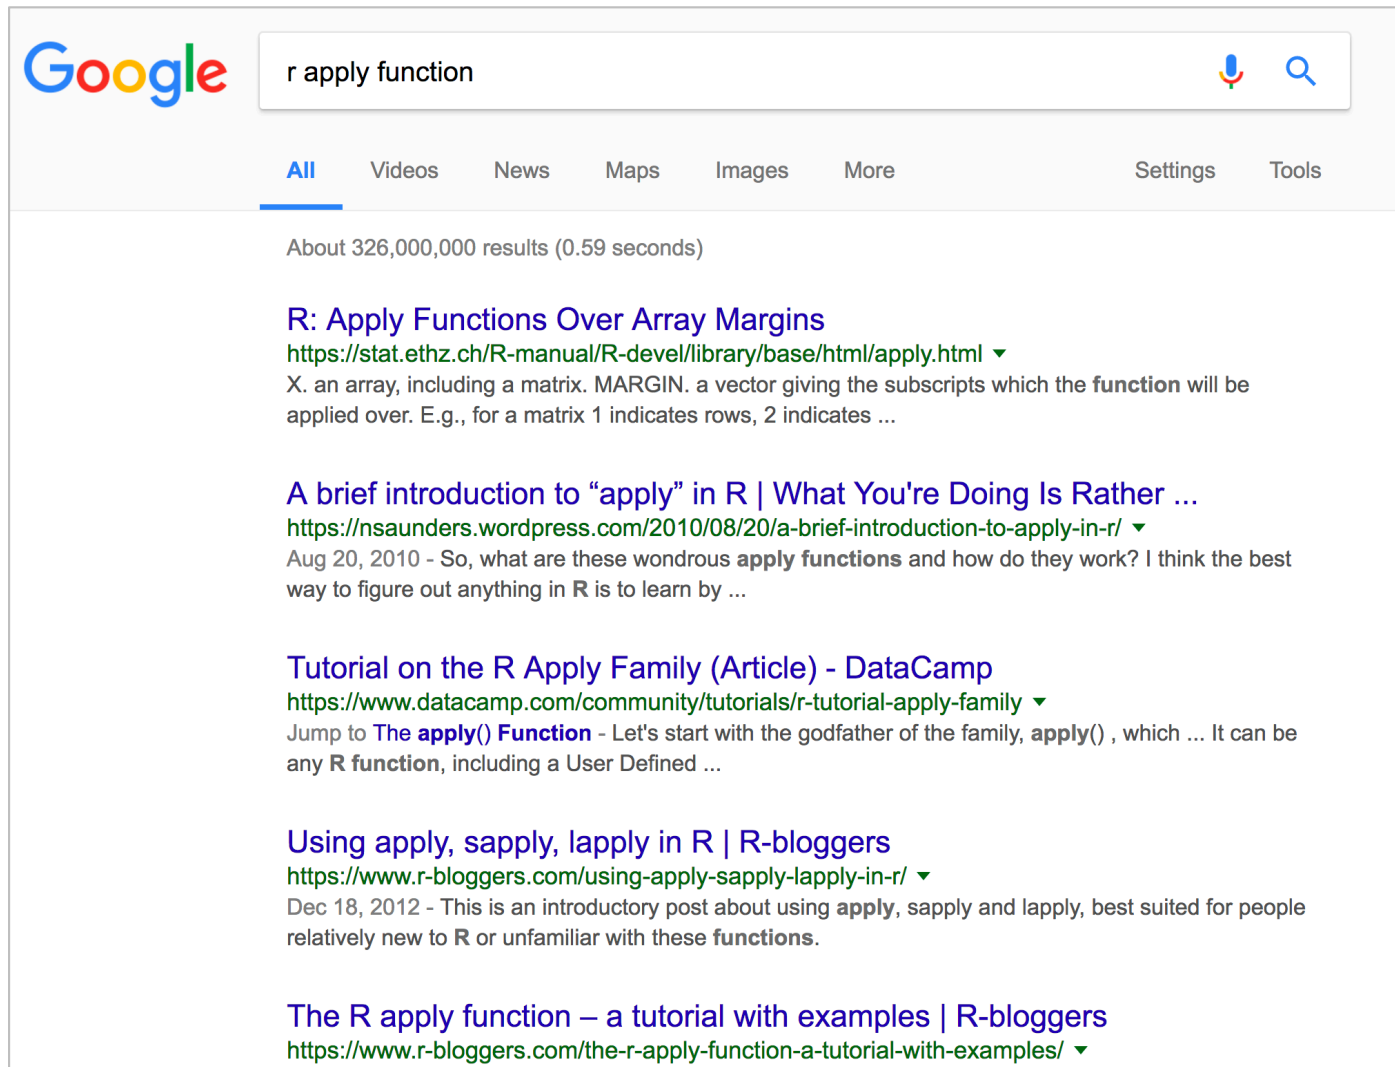
```
apply(X, MARGIN, FUN, ...)
```

```
Arguments:
```

```
X: an array, including a matrix.
```

Tip: type **q** to exit at any time

Finding Help: Google




The screenshot shows a Google search interface with the query "r apply function". The search results are displayed under the "All" tab. The results include:

- R: Apply Functions Over Array Margins**
<https://stat.ethz.ch/R-manual/R-devel/library/base/html/apply.html> ▼
X. an array, including a matrix. MARGIN. a vector giving the subscripts which the **function** will be applied over. E.g., for a matrix 1 indicates rows, 2 indicates ...
- A brief introduction to “apply” in R | What You're Doing Is Rather ...**
<https://nsaunders.wordpress.com/2010/08/20/a-brief-introduction-to-apply-in-r/> ▼
Aug 20, 2010 - So, what are these wondrous **apply functions** and how do they work? I think the best way to figure out anything in **R** is to learn by ...
- Tutorial on the R Apply Family (Article) - DataCamp**
<https://www.datacamp.com/community/tutorials/r-tutorial-apply-family> ▼
Jump to **The apply() Function** - Let's start with the godfather of the family, **apply()** , which ... It can be any **R function**, including a User Defined ...
- Using apply, sapply, lapply in R | R-bloggers**
<https://www.r-bloggers.com/using-apply-sapply-lapply-in-r/> ▼
Dec 18, 2012 - This is an introductory post about using **apply**, **sapply** and **lapply**, best suited for people relatively new to **R** or unfamiliar with these **functions**.
- The R apply function – a tutorial with examples | R-bloggers**
<https://www.r-bloggers.com/the-r-apply-function-a-tutorial-with-examples/> ▼


Choose the ones that make sense to you!

Basic R Data Types

Visual summary of main R data types




vector



matrix

data.frame



list

Atomic Data Types

- In R the base type is a **vector**.
 - An indexed set of values that are all of the same type.
 - [1, 2, 3, 4, 5], ["a", "b", "c" ...]
 - Mixing different types may cause unexpected results
- The type of the entries determines the class of the vector:
 - Integer, numeric, character, logical
- Check the class of a vector
 - class()
- Force to convert to a particular type:
 - as.integer(), as.numeric(), as.character(), as.logical()

Vector Types

```
> a <- 1:100
> class(a)
[1] "integer"
> a <- letters[1:10]
> a
[1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j"
> class(a)
[1] "character"
> a <- rnorm(10) # built-in function, use "?rnorm" for more details
> a
[1] -0.6768672  0.5078512  0.3473145  1.2601923 -0.1431014 -1.0052774
[7]  0.7683541  0.7247683 -0.9169387  1.3272359
> class(a)
[1] "numeric"
> a <- a > 0
> a
[1] FALSE  TRUE  TRUE  TRUE FALSE FALSE  TRUE  TRUE FALSE  TRUE
> class(a)
[1] "logical"
```

Vector Creation

- Concatenation (**c**), colon (**:**), seq, rep

```
> v1 <- c(1, 2, 3)
> v1
[1] 1 2 3
> v2 <- 1:3
> v2
[1] 1 2 3
> v3 <- seq(from=1, to=10, by=2)
> v3
[1] 1 3 5 7 9
> v4 <- rep(1, 4)
> v4
[1] 1 1 1 1
```

Vector Indexing

- Get vector elements: `[]`
 - By numeric index
 - ** Index start with 1, not 0
 - Use positive index to get specific element
 - Use negative to exclude elements
 - By name
 - For named vectors
 - By Boolean index (vector with logical values)
 - Same length
 - Return element with TRUE

```
> v <- c(1, 2, 3)
> v[2]
[1] 2
> v[-2]
[1] 1 3
> v[2:3]
[1] 2 3
> names(v) <- c("a", "b", "c")
> v["b"]
b
2
```

Vector Operations

- Boolean vector, common functions, missing values,

```
> test <- rnorm(5)
> test
[1] 0.2255703 0.3354089 -0.7202021 -0.9523239 -1.5018892
> pos.inx <- test > 0
> pos.inx
[1] TRUE TRUE FALSE FALSE FALSE
> test[pos.inx] # using boolean index to get elements
[1] 0.2255703 0.3354089
> mean(test)
[1] -0.5226872
> sd(test)
[1] 0.7871872
> test <- c(test, NA)
> is.na(test)
[1] FALSE FALSE FALSE FALSE FALSE TRUE
```

Factors

- A factor is a character vector augmented with information to store categorical data: **factor()**
 - Used to specify experimental design

```
> d1 <- c("M", "F", "M", "F", "F", "F")
> d2 <- factor(d1)
> d2
[1] M F M F F F
Levels: F M
> table(d2)
F M
4 2
```

Matrix

- A matrix is table of same type: **matrix()**

```
> mat <- matrix(1:12, nrow=3, byrow=T)
> mat
      [,1] [,2] [,3] [,4]
[1,]    1    2    3    4
[2,]    5    6    7    8
[3,]    9   10   11   12

> mat2 <- matrix(letters[1:8], nrow=2, byrow=F)
> mat2
      [,1] [,2] [,3] [,4]
[1,] "a"  "c"  "e"  "g"
[2,] "b"  "d"  "f"  "h"

> dim(mat2) # size of the matrix
[1] 2 4
```

Matrix Indexing

- Use the row and column positions to select values

```
> mat
      [,1] [,2] [,3] [,4]
[1,]    1    2    3    4
[2,]    5    6    7    8
[3,]    9   10   11   12
> mat[2,3] # single value
[1] 7
> mat[2,]  # row selection
[1] 5 6 7 8
> mat[,3]  # column selection
[1] 3 7 11
```

Note, the comma “,” matters!

Lists

- A collection of vectors or other data objects. It is a very flexible way to organize relevant information: **list()**

```
> x <- list(name="Jeff", height=180, language=c("R", "Java",  
"Perl"))  
> x  
$name  
[1] "Jeff"  
  
$height  
[1] 180  
  
$language  
[1] "R"      "Java"   "Perl"
```

List Indexing

- Elements can be accessed by **name** or by **index** `[` or `[[`

```
> x$name
[1] "Jeff"
> a <- x[1] # [ return a list
> a
$name
[1] "Jeff"
> class(a)
[1] "list"
> b <- x[[1]] # [[ return a vector
> b
[1] "Jeff"
> class(b)
[1] "character"
```

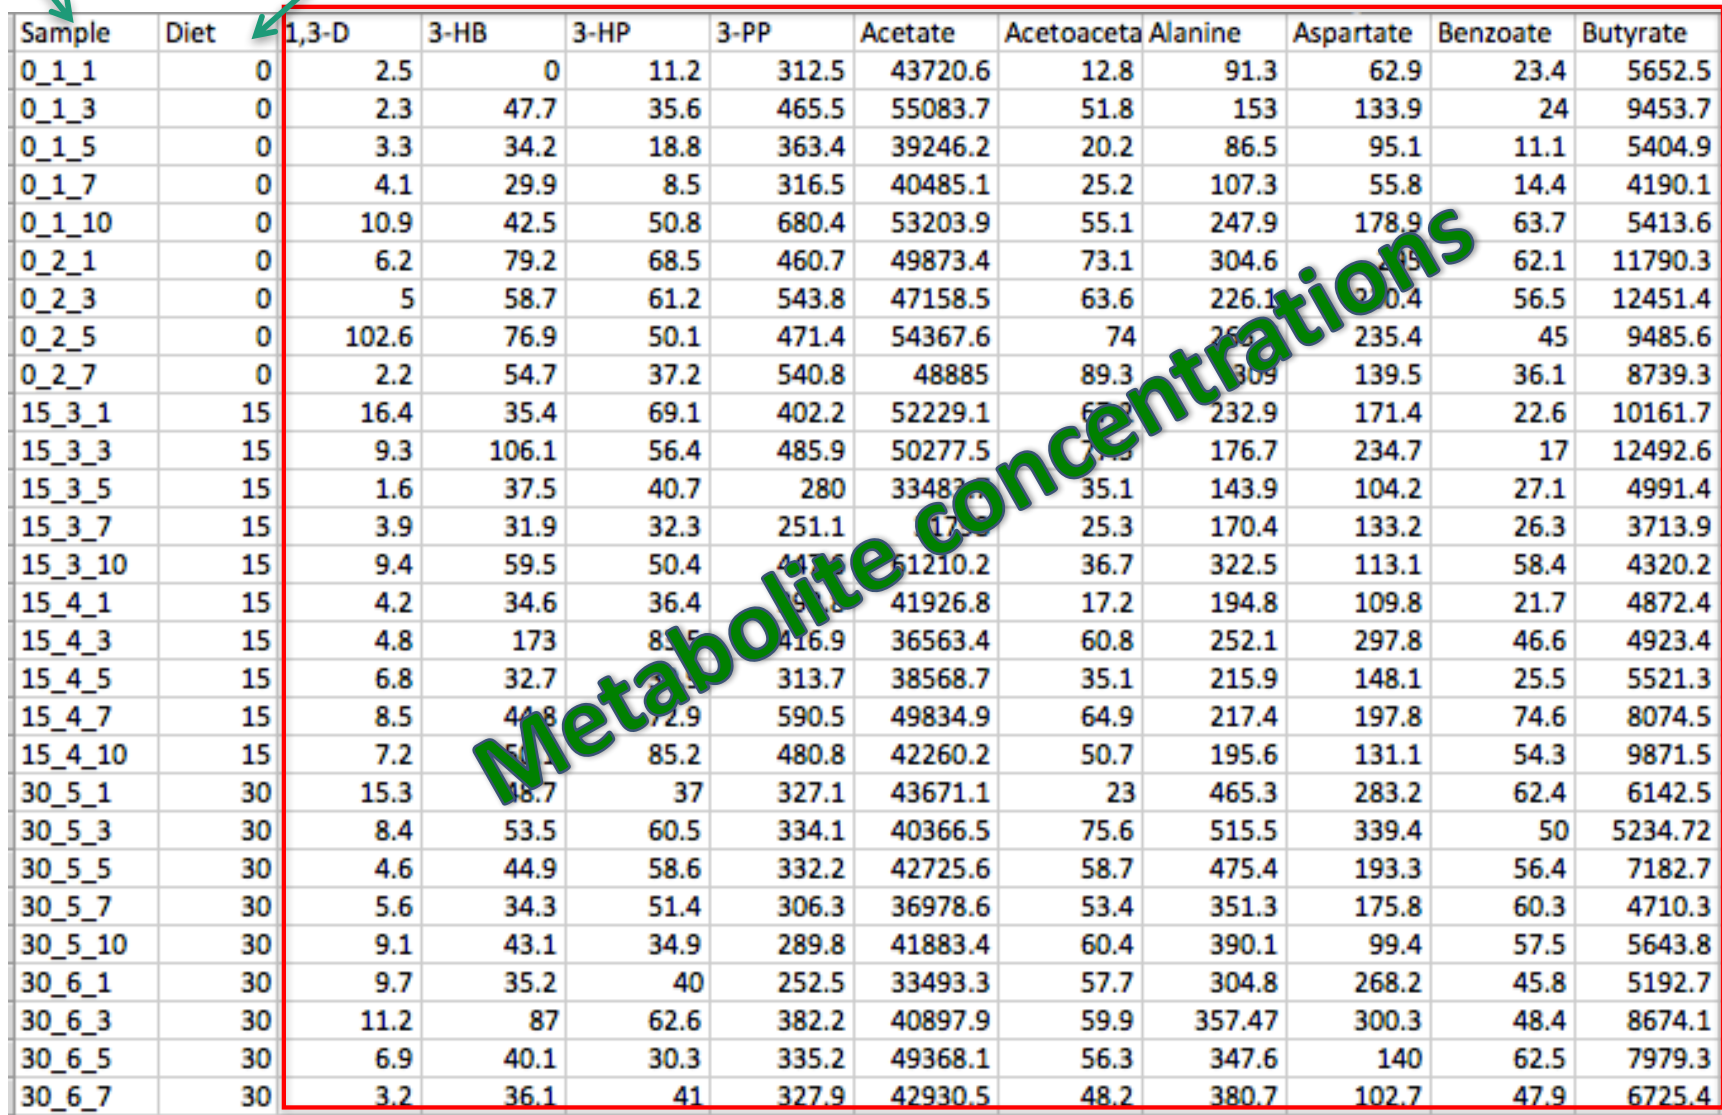
Data Frames

- Data frame is for storing tabular data such as Excels or CSV file: **data.frame()**
- shares many of the properties of matrices and of lists
 - Matrix-like (table) structure
 - Each column is a list of same length

```
> df <-data.frame(  
  ID=c("Mike", "Tom", "Jon", "Mia"),  
  age=c(12, 13, 12, 14),  
  score=c(80, 85, 88, 95) );  
  
> df  
   name age score  
1 Mike  12   80  
2  Tom  13   85  
3  Jon  12   88  
4 Mia  14   95  
  
> df$age  
[1] 12 13 12 14  
  
> df[2,3]  
[1] 85
```

An example data table

Sample IDs Group Labels



Sample	Diet	1,3-D	3-HB	3-HP	3-PP	Acetate	Acetoaceta	Alanine	Aspartate	Benzoate	Butyrate
0_1_1	0	2.5	0	11.2	312.5	43720.6	12.8	91.3	62.9	23.4	5652.5
0_1_3	0	2.3	47.7	35.6	465.5	55083.7	51.8	153	133.9	24	9453.7
0_1_5	0	3.3	34.2	18.8	363.4	39246.2	20.2	86.5	95.1	11.1	5404.9
0_1_7	0	4.1	29.9	8.5	316.5	40485.1	25.2	107.3	55.8	14.4	4190.1
0_1_10	0	10.9	42.5	50.8	680.4	53203.9	55.1	247.9	178.9	63.7	5413.6
0_2_1	0	6.2	79.2	68.5	460.7	49873.4	73.1	304.6	149.5	62.1	11790.3
0_2_3	0	5	58.7	61.2	543.8	47158.5	63.6	226.1	200.4	56.5	12451.4
0_2_5	0	102.6	76.9	50.1	471.4	54367.6	74	266	235.4	45	9485.6
0_2_7	0	2.2	54.7	37.2	540.8	48885	89.3	309	139.5	36.1	8739.3
15_3_1	15	16.4	35.4	69.1	402.2	52229.1	67.2	232.9	171.4	22.6	10161.7
15_3_3	15	9.3	106.1	56.4	485.9	50277.5	74.3	176.7	234.7	17	12492.6
15_3_5	15	1.6	37.5	40.7	280	33483	35.1	143.9	104.2	27.1	4991.4
15_3_7	15	3.9	31.9	32.3	251.1	173.9	25.3	170.4	133.2	26.3	3713.9
15_3_10	15	9.4	59.5	50.4	417.6	51210.2	36.7	322.5	113.1	58.4	4320.2
15_4_1	15	4.2	34.6	36.4	39.8	41926.8	17.2	194.8	109.8	21.7	4872.4
15_4_3	15	4.8	173	83.5	416.9	36563.4	60.8	252.1	297.8	46.6	4923.4
15_4_5	15	6.8	32.7	32.1	313.7	38568.7	35.1	215.9	148.1	25.5	5521.3
15_4_7	15	8.5	44.8	72.9	590.5	49834.9	64.9	217.4	197.8	74.6	8074.5
15_4_10	15	7.2	51.5	85.2	480.8	42260.2	50.7	195.6	131.1	54.3	9871.5
30_5_1	30	15.3	48.7	37	327.1	43671.1	23	465.3	283.2	62.4	6142.5
30_5_3	30	8.4	53.5	60.5	334.1	40366.5	75.6	515.5	339.4	50	5234.72
30_5_5	30	4.6	44.9	58.6	332.2	42725.6	58.7	475.4	193.3	56.4	7182.7
30_5_7	30	5.6	34.3	51.4	306.3	36978.6	53.4	351.3	175.8	60.3	4710.3
30_5_10	30	9.1	43.1	34.9	289.8	41883.4	60.4	390.1	99.4	57.5	5643.8
30_6_1	30	9.7	35.2	40	252.5	33493.3	57.7	304.8	268.2	45.8	5192.7
30_6_3	30	11.2	87	62.6	382.2	40897.9	59.9	357.47	300.3	48.4	8674.1
30_6_5	30	6.9	40.1	30.3	335.2	49368.1	56.3	347.6	140	62.5	7979.3
30_6_7	30	3.2	36.1	41	327.9	42930.5	48.2	380.7	102.7	47.9	6725.4

Data Frames Operation

- Download the “cow_diet.csv” example data set from the MetaboAnalyst site and save it to your current folder
 - http://www.metaboanalyst.ca/resources/data/cow_diet.csv

```
> cow.dat <- read.csv("cow_diet.csv", header=T, row.names=1,
check.names=F, as.is=T)
> class(cow.dat)
[1] "data.frame"
> cow.dat[1:3, 1:4]
```



	Diet	1,3-D	3-HB	3-HP
0_1_1	0	2.5	0.0	11.2
0_1_3	0	2.3	47.7	35.6
0_1_5	0	3.3	34.2	18.8

R Graphics

Box plots

- Visualize distributions of multiple variables

```
> dim(cow.dat) # get the row and column numbers of the data
[1] 39 48
> boxplot(cow.dat[,2:6])
# set the labels orthognal to axis (las=2)
> boxplot(log(cow.dat[,2:6]), las=2, horizontal=T, col=2:6)
```



Question #1

- When you take log on the data, you see the warning message:

```
> boxplot(log(cow.dat[,2:6]), las=2, horizontal=T, col=2:6)
Warning message:
In bplt(at[i], wid = width[i], stats = z$stats[, i], out = z$out[z$group ==  :
  Outlier (-Inf) in boxplot 2 is not drawn
> log(0)
[1] -Inf
```

- How to get rid of the warning?

Solution to Q1



```
# make a copy of the original data, exclude group labels (1st col)
> conc.dat <- cow.dat[,-1]
# get indices for all positive values
> pos.inx <- conc.dat > 0
# get minimal concentration
> min(conc.dat[pos.inx])
[1] 1.6
# replace zero/negative with 1/5 of minimal concentration
> conc.dat[!pos.inx] <- 1.6/5;
> boxplot(log(conc.dat[,1:5]), las=2, horizontal=T, col=2:6)
```

Note, *pos.inx* is a matrix of Boolean values with TRUE for positive values. Apply NOT operator (!) will obtain a matrix with TRUE for non-positive values

Histogram

- Visualize the distributions of numerical values

```
# convert to a numerical vector  
> concs.all <- as.numeric(as.matrix(conc.dat));  
> dim(concs.all)  
NULL  
> hist(concs.all)  
> hist(concs.all, xlab="Concentrations", main="")  
> hist(log(concs.all), xlab="Concentrations", col="blue",  
breaks=30, main="")
```




Question #2

- Based on the histogram, the log transformed data between 1.5 to 8.5 seems to be normally distributed, how do you verify that?

Solution to Q2



```
> norm.inx <- concs.all > 1.5 & concs.all < 8.5  
> norm.dat <- concs.all[norm.inx];  
> qqnorm(norm.dat)  
> qqline(norm.dat);
```



Scatter plot

- Visualize the correlations between two variables

```
> plot(cow.dat[,4:5])  
> plot(cow.dat[,4:5], type='n')  
> grid()  
> points(cow.dat[,4:5], pch=19, col='darkblue')
```



Functions & Scripts

Built-in Functions

- Functions are a set of commands that work together to perform a given task
- Arguments are parameters you provide to the function for processing

```
> sum(c(1,2,3))  
[1] 6  
  
> mean(c(1,2,3,4,5))  
[1] 3  
  
# get average concentraion of each compound  
> apply(conc.dat, 2, mean)  
  
# get total concentration values of each sample  
> apply(conc.dat, 1, sum)
```

Writing Functions: *function_name (args)*

- Functions are a set of commands that work together to perform a given task
- Let's convert our previous solution to a function

```
# Replace zero and negative values with 1/5 of min positive values
# input: a numerical matrix containing zero or negative values
# output: matrix with all positive values
ReplaceNonpositives <- function(data) {
  pos.inx <- data > 0;
  min.val <- min(data[pos.inx]);
  data[!pos.inx] <- min.val/5;
  return(data);
}
```


Write R scripts

An R script is just a plain text file with R commands & functions in it. You can prepare a script in any text editor.

1. Save the function in a text file named “ReplaceNonpositives.R” under current directory
 - You can choose other informative names
 - Multiple functions can be saved in the same file
2. To use the function:

```
> source("ReplaceNonpositives.R")  
> conc.dat2 <- ReplaceNonpositives(conc.dat);
```

Libraries & Packages

- One important reason R has become so popular is the vast array of packages available at the **CRAN** and **Bioconductor** repositories

- Install & load a CRAN package

```
> install.packages("ggplot2")  
> library(ggplot2)
```


- Install a Bioconductor package

```
> source("http://bioconductor.org/biocLite.R")  
> > biocLite("multtest")
```

Finding help

- Google your questions!
- R website <http://www.r-project.org>
 - Documentation
 - Mailing-list R-help
- Useful links
 - <http://www.r-bloggers.com>
 - <http://stackoverflow.com/>
 - <http://www.biostars.org>

Statistical Bioinformatics (BINF 531)



Sizes are proportional to the actual difficulties

Any questions?